

DATADYNAMOS SYSTEM ARCHITECTURE SPECIFICATION

Comprehensive End-to-End Pipeline, Multi-Engine OCR Orchestration & Rule Engine Blueprint


1. System Architectural Overview & Design Philosophy

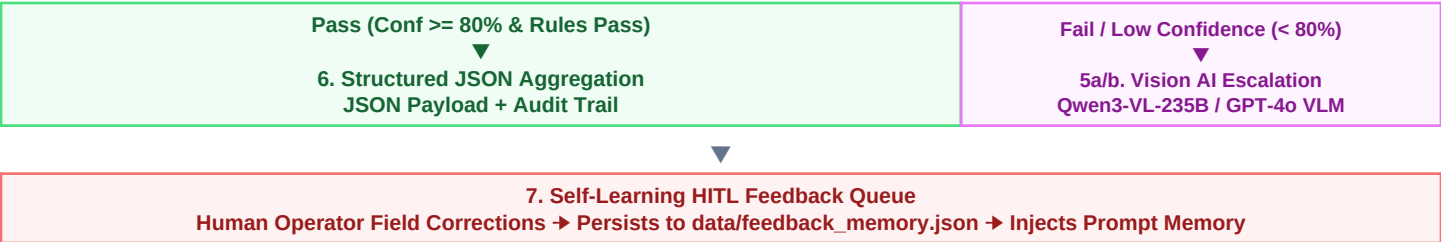
DataDynamos is an enterprise-grade, zero-retraining platform engineered to process healthcare forms (**CMS-1500**, **UB-04**, **Unstructured Medical Claims**), Commercial Invoices, and Legal Contracts at a target scale of **100 Million pages per year** at **\$0.000375 / page average cost** with **93.5% Straight-Through Processing (STP)**.

- CPU-First Cost Optimization:** 97% of machine-printed forms run on zero-cost, C++ CPU-bound OCR engines (PaddleOCR/PyTesseract at \$0.0002/pg), reserving expensive Vision-Language Models (Qwen3-VL-235B at \$0.0030/pg) only for low-confidence noisy scans.
- Deterministic Compliance Primacy:** Hard Python rule engines (NPI Luhn 10-digit checksums, ICD-10-CM format audit, UB-04 revenue math balancing) take precedence over LLM output, preventing AI hallucination in financial/medical claims.
- Zero-Retraining Prompt Memory HITL Loop:** Inline human corrections captured in the React inspector persist to `data/feedback_memory.json` and inject learned prompt memory into future LLM extractions without model weight retraining.
- Spatial Bounding-Box JSON Feeding:** Raw OCR outputs are serialized into rich spatial JSON (bounding box coordinates + Markdown tables) fed into OpenRouter LLMs via the LangExtract framework for high-precision field extraction.

2. End-to-End System Architecture Flowchart

The following multi-tier processing flowchart illustrates document classification, dynamic OpenCV preprocessing, template/anchor vs layout OCR parsing, deterministic rule validation, VLM escalation routing, and the self-learning HITL feedback memory queue:





3. Detailed Stage-by-Stage Implementation Matrix

Stage	Pipeline Layer & Module	Input → Core Operation → Output Artifact	Latency & Cost
Stage 1	Pre-Scan Quality Engine app/pipeline/prescan.py	Incoming PDF/Scan → OpenCV Deskew (±25°), DPI Normalization (200 DPI), Blur/Contrast Check → Cleaned BGR Raster & Quality Report	45 ms \$0.0000
Stage 2	Format AI Classifier app/pipeline/classifier.py	Raster Page → Layout & Keyword Heuristics → Tier A (CMS-1500 Single), Tier B (CMS-1500 Multi), Tier C (UB-04), Tier D (Unstructured)	15 ms \$0.0000
Stage 3	Multi-Engine OCR Orchestrator app/pipeline/ocr/*	Clean Raster → PaddleOCR CPU / PyTesseract / Docling (Tables) / Qwen3-VL Escalation (Noise <80%) → Block OCR JSON + `bbox` Coordinates	320 ms \$0.0002
Stage 4	Spatial LLM Field Structurer app/pipeline/structuring.py	Block OCR JSON → LangExtract Spatial JSON Serialization → OpenRouter LLMs (DeepSeek-v4 / GPT-4o) → Structured JSON + Grounding Map	680 ms \$0.00012
Stage 5	Deterministic Rule Audit Engine app/rules/*	Structured JSON → NPI Luhn Checksum, ICD-10 Audit, CPT Code Check, UB-04 Revenue Line Math Balance → Audit Check Results Array	10 ms \$0.0000
Stage 6	Autonomous Decision Agent app/pipeline/agent.py	Audit Checks + LLM Reasoning → Reconcile Code Rules & LLM Judgment → Verdict: `approve` `needs_review` `flag`	350 ms \$0.000055
Stage 7	Self-Learning HITL Feedback Loop app/routes/documents.py	Human Operator Corrections → Persist to `data/feedback_memory.json` → Inject Dynamic Exemplars in Structuring Prompts	Async \$0.0000

4. Multi-Engine OCR Orchestration & Dynamic Escalation Logic

To achieve sub-\$0.0004 per page costs across 100M pages, DataDynamos implements a tiered OCR orchestration routing architecture:

- **PaddleOCR (PP-OCRv4 CPU) (\$0.0002 / page):** Primary default engine for standard machine-printed grid forms (CMS-1500). Executes in C++ on CPU with ~1.2s per page latency and zero commercial API costs.
- **PyTesseract (v5.3) (\$0.0001 / page):** Ultra-fast fallback CPU engine for standard monospaced text lines and fixed-layout institutional forms (UB-04).
- **Docling Deep Layout Parser (\$0.0003 / page):** Specialized layout analysis model for multi-column documents, complex tables, and Markdown tabular extraction.
- **Qwen3-VL-235B Vision AI (\$0.0030 / page):** Multimodal Vision-Language Model invoked over OpenRouter *conditionally* when OCR confidence drops below 80% or severe physical skew/staining is detected.

5. Deterministic Rule Audit Engine & ANSI Healthcare Compliance

Rule Identifier	Target Document Tier	Validation & Compliance Audit Logic	Severity & Action
billing_npi_luhn [ANSI B1]	CMS-1500 / UB-04	Validates 10-digit National Provider Identifier using US prefix `80840` and ANSI Luhn checksum algorithm.	Hard Failure (Forces `flag`)
icd10_cm_format [ANSI C1]	CMS-1500 / UB-04	Verifies ICD-10-CM diagnosis code syntax (1 letter + 2 digits + optional decimal + 1-4 characters).	Hard Failure (Forces `flag`)
revenue_charges_balance [ANSI D2]	UB-04 Institutional	Verifies that the sum of individual revenue line item charges (Box 47) equals the total billed amount.	Hard Failure (Forces `flag`)
patient_identity_match [ANSI A1]	All Healthcare Claims	Ensures Patient Name, DOB, and Policy ID exist and match cross-field references.	Review Gate (Caps at `needs_review`)
extraction_confidence	Cross-Cutting (All)	Ensures average extraction confidence across all fields exceeds 80%.	Review Gate (Caps at `needs_review`)

6. Technology Stack & Component Mapping

Layer	Technology / Framework	Role & Technical Implementation
Backend API	FastAPI + Uvicorn	Asynchronous REST API with auto-generated OpenAPI docs (`:8000`).
Package Manager	Astral `uv`	Lightning-fast Python dependency locking and virtualenv sync.
Computer Vision	OpenCV 4.x + PyMuPDF	Deskew, binarization, DPI scaling, and raster rendering.
OCR Engines	PaddleOCR + Tesseract + Docling	Multi-engine CPU OCR and layout parsing orchestration.
LLM Structuring	LangExtract Framework	Structured JSON extraction over OpenRouter LLMs (DeepSeek-v4, GPT-4o).
Frontend SPA	React 18 + Vite + TypeScript	Glassmorphism dashboard with split inspector and bounding box highlights.
UI Components	TailwindCSS v4 + Lucide Icons	Modern responsive UI design system with interactive inspector tabs.
Reporting & Export	ReportLab + Pandas + OpenPyXL	Automated PDF specification and Excel benchmark generation.

7. Enterprise Deployment & Scalability Architecture

To support **100 Million pages per year** (equivalent to 70.4 pages per second across a 100-worker cluster):

- **Stateless Worker Nodes:** FastAPI backend workers execute statelessly in Docker containers, allowing horizontal auto-scaling on Kubernetes.
- **Docker Compose Infrastructure:** Unified multi-stage containerization orchestrating backend API (`:8000`) and Vite frontend (`:5173`).
- **Zero-GPU Operational Flexibility:** Standard claim forms run 100% on CPU worker nodes, eliminating expensive GPU cloud infrastructure dependencies.